

Hand Gesture Recognition using Deep Learning

1. Introduction

This project develops a deep learning system to recognize dynamic hand gestures from video sequences, facilitating touchless human-computer interaction. The system uses Transfer Learning with pre-trained CNNs and a Time-Distributed architecture to handle temporal video data.

2. Dataset & Preprocessing

- **Source:** [Kaggle Hand Gesture Recognition Dataset](#).
- **Classes:** Left_Swipe, Right_Swipe, Stop, Thumbs_Down, Thumbs_Up.
- **Preprocessing:**
 - Frame Extraction: Videos decomposed into frames.(30)
 - Resizing: Standardized to 64x64 pixels.
 - Normalization: Pixels scaled to [0, 1].
 - Augmentation: Horizontal flipping applied to balance classes.
 - Serialization: Saved as .npy files with metadata management.

3. Methodology


Three models were implemented using Transfer Learning. All follow a consistent architecture designed for video processing.

All models process video batches of shape (30 Frames, 64, 64, 3)

1. Model 1:VGG19


- **Base:** Custom VGG19-style CNN (built with Conv2D + BatchNorm + ReLU blocks).
- **Focus:** Classic deep CNN feature learning (stacked 3×3 convolutions + max pooling).
- **Training:** Trained from scratch (all layers trainable; no pretrained backbone).
- **Overall structure**
 - Input: (224×224×3)
 - Backbone (convolutional feature extractor):
 - Block 1: [Conv3×3(64) → BN → ReLU] ×2 → MaxPool(/2)
 - Block 2: [Conv3×3(128) → BN → ReLU] ×2 → MaxPool(/2)
 - Block 3: [Conv3×3(256) → BN → ReLU] ×4 → MaxPool(/2)
 - Block 4: [Conv3×3(512) → BN → ReLU] ×4 → MaxPool(/2)
 - Block 5: [Conv3×3(512) → BN → ReLU] ×4 → MaxPool(/2)
- **Top / head:**
 - Flatten
 - Dense(4096, use_bias=False) → BN → ReLU → Dropout(0.5)
 - Dense(4096, use_bias=False) → BN → ReLU → Dropout(0.5)
 - Dense(num_classes=5, activation="softmax")

- **Trainability pattern**
 - Phase 1: all layers trainable (trained end-to-end from scratch)
 - Phase 2: none (no freezing / fine-tuning stage in the notebook)
- **Diagram:**




2. Model 2: ResNet50

- **Base:** ResNet50.
- **Focus:** Deep network training via residual connections.
- **Training:** Frozen base (Feature Extraction only).
- **Overall structure**
 - Input: $(256, 256, 3)$
 - Backbone: ResNet50(include_top=False, weights="imagenet", pooling="avg")
- **Top / head:**
 - Dense(num_classes, activation="softmax")
 - Output: class probabilities
- **Trainability pattern**
 - Phase 1: backbone frozen (base.trainable = False)
 - Phase 2: backbone partially fine-tuned (only last UNFREEZE_LAST_LAYERS layers trainable; earlier layers frozen)
- **Diagram:**



3. Model 3: InceptionV1


- **Base:** InceptionV1 (TensorFlow Hub).
- **Focus:** Multi-scale feature capture using Inception modules.
- **Training:** Frozen base (Feature Extraction only).
- **Overall structure**
 - Input: (30' 64' 64' 3) (*30-frame clip*)
 - Per-frame backbone (frozen): TF-Hub
 - TimeDistributed (Lambda(feature_extractor)) over the 30 frames
 - GlobalAveragePooling1D()
- **Top / head:**
 - Dense(256, activation="relu")
 - Dropout(0.3)
 - Dense(num_classes, activation="softmax")
 - Output: class probabilities
- **Trainability pattern**
 - Backbone stays frozen (trainable=False); only the head trains.
- **Diagram:**



4. Model 4: MobileNetV2

- **Base:** MobileNetV2 (ImageNet weights).
- **Focus:** Lightweight efficiency for mobile devices.
- **Model structure**

- Input: video clip : (30, 64, 64, 3)
- Per-frame preprocessing:
 - TimeDistributed (Resizing(IMG_SIZE=128, 128))
 - TimeDistributed (Rescaling(2.0, offset=-1.0)) (*maps [0,1] → [-1,1]*)
- Applied as TimeDistributed(base_model) to each frame
- Temporal aggregation: GlobalAveragePooling1D() across the 30 frame-vectors
- Output: class probabilities
- **Top / head:**
 - Dense(256, activation="relu")
 - Dropout (0.3)
 - Dense (num_classes, activation="softmax")
- **Trainability / fine-tuning pattern**
 - **Phase 1:** base_model.trainable = False (backbone fully frozen)
 - **Phase 2:** base_model.trainable = True
- **Diagram:**



4. Evaluation Metrics:

- **Accuracy:** Overall correct classification percentage.
- **Confusion Matrix:** Visualization of specific misclassifications.
- **Classification Report:** Precision, Recall, and F1-Score.
- **ROC & AUC Curves:** Evaluation of discrimination capability.

5.Comparison

<i>Model Architecture</i>	<i>Training Strategy</i>	<i>Input Resolution</i>	<i>Val/Test Accuracy</i>	<i>Model Status</i>
ResNet50	Transfer Learning (ImageNet)	256 x 256	~80 - 81%	Highest Accuracy
MobileNetV2	Transfer Learning (ImageNet)	128 x 128	78.26%	Most Efficient
Inception V1	Transfer Learning (TF Hub)	224 x 224	75.65%	Moderate Performance
VGG19	Training from Scratch	224 x 224	20.00%	Failed to Converge

5.1. Conclusion & Selection

1. **Deployment Choice:** MobileNetV2 is selected as the production model. Its lightweight architecture ensures low latency during real-time inference while maintaining high classification reliability.
2. **Benchmark Choice:** ResNet50 serves as the accuracy benchmark, demonstrating the upper limit of performance extractable from the current dataset using standard transfer learning techniques.

6. References

1. VGG-19 model:

- *Paper Title:* ["Very Deep Convolutional Networks for Large-Scale Image Recognition"](#)
- *Authors:* Karen Simonyan, Andrew Zisserman (University of Oxford).
- *Published In:* International Conference on Learning Representations (ICLR), 2015.

2. ResNet50:

- *Paper:* ["Deep Residual Learning for Image Recognition"](#)
- *Authors:* Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun (2016).
- *Code Documentation:* [Keras ResNet50 API Docs](#)

3. MobileNet:

- *Paper:* ["MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications"](#)
- *Authors:* Andrew G. Howard et al. (2017).
- *Code Documentation:* [Keras MobileNet API Docs](#)

4. Inception V1 (GoogLeNet):

- *Paper:* ["Going Deeper with Convolutions"](#)

- *Authors:* Christian Szegedy et al. (2015).
- *Code Documentation:* [*Keras MobileNet API Docs*](#)

5. **Dataset Source:** [Kaggle: Hand Gesture Recognition Dataset.](#)